

Introduction to Linear Regression and ANOVA with R

Laura Gutierrez Funderburk

Simon Fraser University

August 2019



SIMON FRASER
UNIVERSITY

- 1 Introductions
- 2 Linear Regression (In a Nutshell)
- 3 ANOVA (In a Nutshell)
- 4 Hands on activities

About Me

- Mathematics Major, SFU
- DevOps Engineer for iReceptor, SFU



SIMON FRASER
UNIVERSITY

When do I use Linear Regression

- When I want to determine whether there is a relationship between two or more variables
- If there is a relationship, I want a model that will help me predict new values from existing data
- Help determine how much I can trust my model



SIMON FRASER
UNIVERSITY

When do I use ANOVA

- To determine if there are statistically significant differences between two or more groups
- Can come in handy when your data involves control and treatment groups (health sciences, biology for example)



Tell me about you

- Name
- What you do
- One challenge you face in your work
- What you hope to learn from this workshop



SIMON FRASER
UNIVERSITY

What is Linear Regression

The purpose of linear regression is to model a continuous variable Y as a mathematical function of one or more X variable(s).

We can use this regression model to predict the Y when only the X is known.

This mathematical equation can be generalized as follows:

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_N X_N + \epsilon$$

where, β_i are the coefficients associated to each X_i and ϵ is the error term - the part that the model cannot explain.

What kinds of models are we going to explore today

- One example of linear (single variable)

$$Y = \beta_0 + \beta_1 X_1 + \epsilon$$

- One example of linear (multivariate)

$$Y = \beta_0 + \beta_1 X_1 + \beta_2 X_2 + \dots + \beta_N X_N + \epsilon$$

The logo of Simon Fraser University, consisting of the letters "SFU" in white on a red rectangular background.

SIMON FRASER
UNIVERSITY

What kinds of models are we going to explore today

- One example of polynomial (quadratic)

$$Y = \beta_0 + \beta_1 X_1^2 + \epsilon$$



SIMON FRASER
UNIVERSITY

What is a one-way ANOVA test

The one-way analysis of variance (ANOVA) is an extension of independent two-samples t-test for comparing means in two or more groups

Data is organized into several groups based on a factor variable



ANOVA test hypotheses:

Null hypothesis H_0 : the means of the different groups are the same

Alternative hypothesis H_A : At least one sample mean is not equal to the others



ANOVA assumptions:

- The observations are obtained independently and randomly from the population defined by the factor levels
- The data of each factor level are normally distributed
- These normal populations have a common variance



What we are going to do today

- Visualize data sets and explore linear, multivariate and polynomial models
- Use R to determine the model that fits the data best
- Use ANOVA to determine if there are statistically significant differences between two or more groups



SIMON FRASER
UNIVERSITY

Time to practice

Open RStudio

If you do not have RStudio, go to <https://sfu.syzygy.ca/>



SIMON FRASER
UNIVERSITY